



Basics of Deep Neural Networks

Understanding Visual Appearance Through Deep Learning
March 4th, 2020

Preethi Jyothi (CSE)

What is deep learning?

“Deep learning allows computational models that are composed of multiple processing layers to learn representations of data with multiple levels of abstraction.”

“Representation learning is a set of methods that allows a machine to be fed with raw data and to automatically discover the representations needed for detection or classification. Deep-learning methods are representation-learning methods with multiple levels of representation, obtained by composing simple but non-linear modules that each transform the representation at one level (starting with the raw input) into a representation at a higher, slightly more abstract level.”

LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436

History of (Deep) Neural Networks

- McCulloch-Pitts Neuron Model (1943)
- Perceptrons (1957)
- Backpropagation (1960)
- Backpropagation for neural networks (1986)
- Convolutional neural networks (1989)
-
-
- Deep learning for speech recognition (2009)
- AlexNet (2012)
- Generative Adversarial Networks (GANs) (2014)
- AlphaGo (2016)

Why the resurgence?

- McCulloch-Pitts Neuron Model (1943)
- Perceptrons (1957)
- Backpropagation (1960)
- Backpropagation for neural networks (1986)
- Convolutional neural networks (1989)
- ⋮
- Deep learning for speech recognition (2009)
- AlexNet (2012)
- Generative Adversarial Networks (GANs) (2014)
- AlphaGo (2016)

Vast amounts of data

+

Specialized hardware,
Graphics Processing Units (GPUs)


+

Improved optimization techniques
and new model variants/libraries/toolkits



Nuts and Bolts of Deep Neural Networks

Feed-forward Neural Network

Single Neuron




Single neuron



Feed-forward Neural Network


Parameterized Model



$$a_3 = g(w_{13} \cdot a_1 + w_{23} \cdot a_2 + b_3)$$

Feed-forward Neural Network

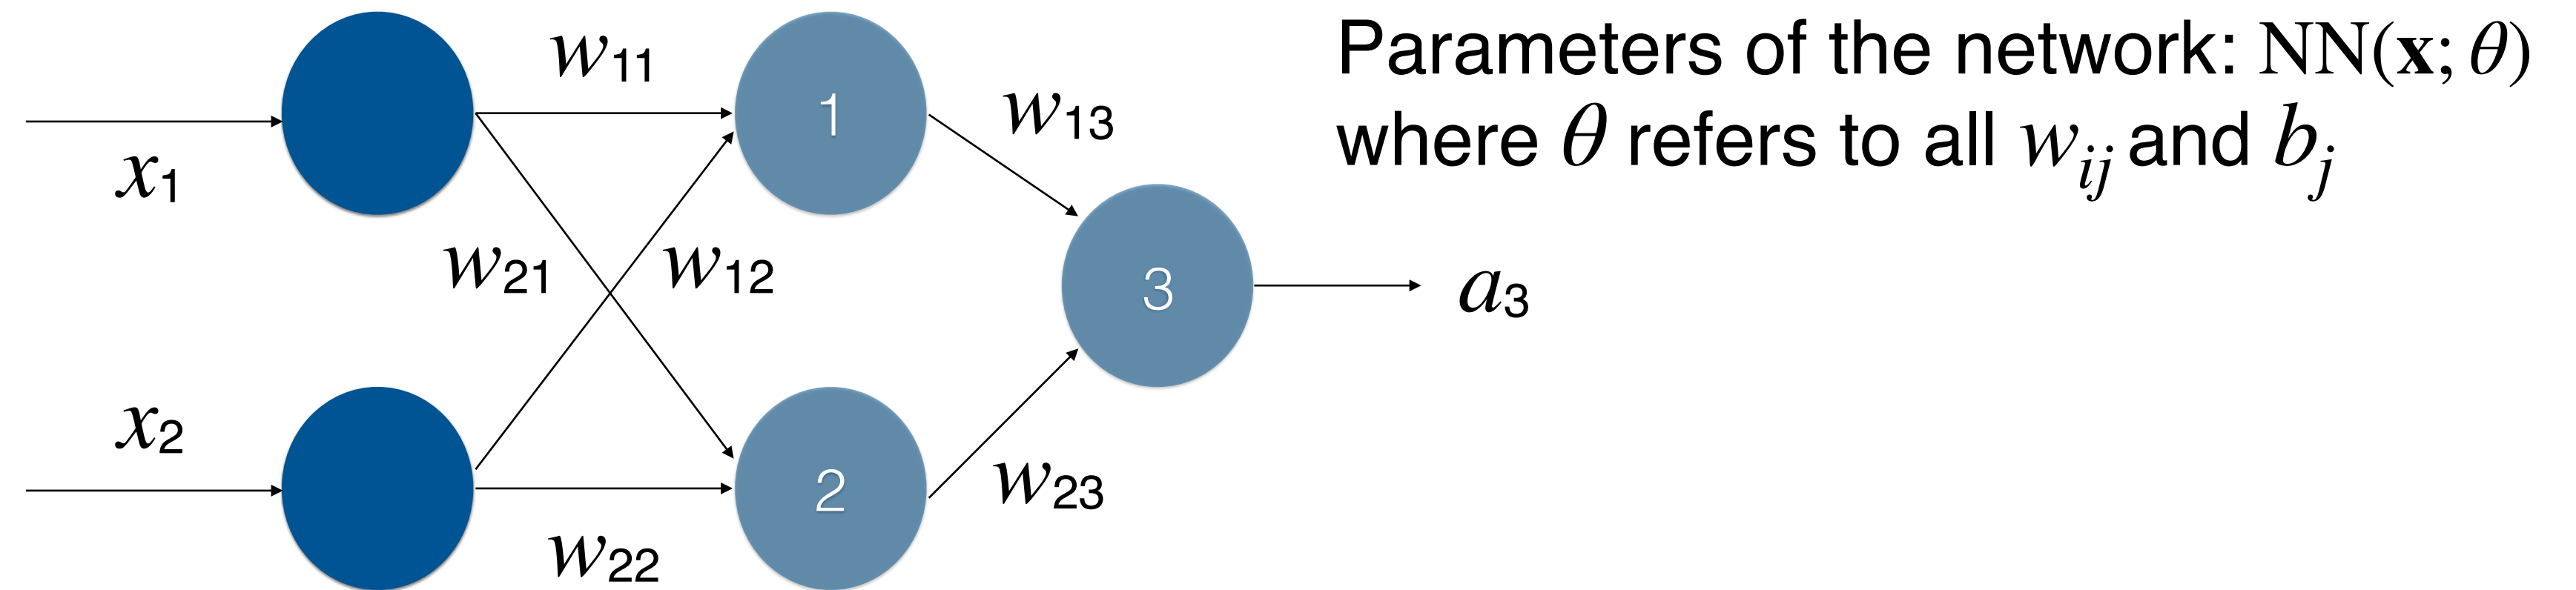
Parameterized Model



$$\begin{aligned} a_3 &= g(w_{13} \cdot a_1 + w_{23} \cdot a_2 + b_3) \\ &= g(w_{13} \cdot (g(w_{11} \cdot x_1 + w_{21} \cdot x_2 + b_1)) \\ &\quad + \dots \end{aligned}$$

Feed-forward Neural Network

Parameterized Model



$$\begin{aligned} a_3 &= g(w_{13} \cdot a_1 + w_{23} \cdot a_2 + b_3) \\ &= g(w_{13} \cdot (g(w_{11} \cdot x_1 + w_{21} \cdot x_2 + b_1)) \\ &\quad + w_{23} \cdot (g(w_{12} \cdot x_1 + w_{22} \cdot x_2 + b_2)) + b_3) \end{aligned}$$


Compact matrix notation: Input $\mathbf{x} = [x_1, x_2]$ is written as a 2-dimensional vector and the layer above it is a 2-dimensional vector \mathbf{h} , a fully-connected layer is associated with:

$$\mathbf{h} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

where w_{ij} in \mathbf{W} is the weight of the connection between i^{th} neuron in the input row and j^{th} neuron in the first hidden layer and \mathbf{b} is the bias vector.

Feed-forward Neural Network

Parameterized Model



$$\begin{aligned} a_3 &= g(w_{13} \cdot a_1 + w_{23} \cdot a_2 + b_3) \\ &= g(w_{13} \cdot (g(w_{11} \cdot x_1 + w_{21} \cdot x_2 + b_1)) \\ &\quad + w_{23} \cdot (g(w_{12} \cdot x_1 + w_{22} \cdot x_2 + b_2)) + b_3) \end{aligned}$$

The simplest neural network is the perceptron:


$$\text{Perceptron}(\mathbf{x}) = \mathbf{x}\mathbf{W} + \mathbf{b}$$

A 1-layer feedforward neural network (multi-layer perceptron) has the form:

$$\text{MLP}(\mathbf{x}) = g(\mathbf{x}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2$$

Common Activation Functions (g)


Sigmoid: $\sigma(x) = 1/(1 + e^{-x})$



Common Activation Functions (g)

Sigmoid: $\sigma(x) = 1/(1 + e^{-x})$

Hyperbolic tangent (tanh): $\tanh(x) = (e^{2x} - 1)/(e^{2x} + 1)$




Common Activation Functions (g)

Sigmoid: $\sigma(x) = 1/(1 + e^{-x})$

Hyperbolic tangent (tanh): $\tanh(x) = (e^{2x} - 1)/(e^{2x} + 1)$

Rectified Linear Unit (ReLU): $\text{ReLU}(x) = \max(0, x)$



Training Neural Networks

Optimization Problem

- To train a neural network, define a loss function $L(y, \tilde{y})$:
a function of the true output y and the predicted output \tilde{y}
- $L(y, \tilde{y})$ assigns a non-negative numerical score to the neural network's output, \tilde{y}
- The parameters of the network are set to minimise L over the training examples (i.e. a sum of losses over different training samples)
- L is typically minimised using a gradient-based method

Stochastic Gradient Descent (SGD)

SGD Algorithm

Inputs: $\text{NN}(x; \theta)$, Training examples, $x_1 \dots x_n$; outputs, $y_1 \dots y_n$ and Loss function L

Randomly initialize θ

do until **stopping criterion**

 Pick a training example $\{x_i, y_i\}$

 Compute the loss $L(\text{NN}(x_i; \theta), y_i)$

 Compute gradient of L , $\nabla_{\theta} L$ with respect to θ

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L$$

Weight
Update Rule

done

Learning
Rate

Return: θ

Mini-batch Gradient Descent (GD)

Mini-batch GD Algorithm

Inputs: $\text{NN}(x; \theta)$, Training examples, $x_1 \dots x_n$; outputs, $y_1 \dots y_n$ and Loss function L

Randomly initialize θ

do until **stopping criterion**

 Randomly sample a batch of training examples $\{x_i, y_i\}_{i=1}^b$

 (where the batch size, b , is a hyperparameter)

 Compute gradient of L over the batch, $\nabla_{\theta} L$ with respect to θ

$\theta \leftarrow \theta - \eta \nabla_{\theta} L$

done

Return: θ

Loss Function

Overall loss function, $J(\theta)$, measures the total loss over the entire training set:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N L(\text{NN}(\mathbf{x}_i; \theta), y_i)$$

Cross-entropy loss is one of the most popular classification-based loss functions. Assuming $\text{NN}(\mathbf{x}_i; \theta)$ returns a probability, binary cross-entropy can be defined as:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \log (\text{NN}(\mathbf{x}_i; \theta)) + (1 - y_i) \log (1 - \text{NN}(\mathbf{x}_i; \theta))$$

Training a Neural Network

Define the Loss function to be minimised as a node L

Goal: Learn weights for the neural network which minimise L

Gradient Descent: Find $\partial L / \partial w$ for every weight w , and update it as
 $w \leftarrow w - \eta \partial L / \partial w$

How do we efficiently compute $\partial L / \partial w$ for all w ?

Will compute $\partial L / \partial u$ for every node u in the network!

$$\partial L / \partial w = \partial L / \partial u \cdot \partial u / \partial w \text{ where } u \text{ is the node which uses } w$$

Training a Neural Network

New goal: compute $\partial L / \partial u$ for every node u in the network

Simple algorithm: Backpropagation

Key fact: Chain rule of differentiation

If L can be written as a function of variables v_1, \dots, v_n , which in turn depend (partially) on another variable u , then


$$\partial L / \partial u = \sum_i \partial L / \partial v_i \cdot \partial v_i / \partial u$$

Backpropagation

If L can be written as a function of variables v_1, \dots, v_n , which in turn depend (partially) on another variable u , then

$$\partial L / \partial u = \sum_i \partial L / \partial v_i \cdot \partial v_i / \partial u$$

Consider v_1, \dots, v_n as the layer above u , $\Gamma(u)$



Then, the chain rule gives

$$\partial L / \partial u = \sum_{v \in \Gamma(u)} \partial L / \partial v \cdot \partial v / \partial u$$

Backpropagation

$$\partial L / \partial u = \sum_{v \in \Gamma(u)} \partial L / \partial v \cdot \partial v / \partial u$$

Backpropagation

Base case: $\partial L / \partial L = 1$

For each u (top to bottom):

For each $v \in \Gamma(u)$:

Inductively, have computed $\partial L / \partial v$

Directly compute $\partial v / \partial u$


Compute $\partial L / \partial u$

Compute $\partial L / \partial w$

where $\partial L / \partial w = \partial L / \partial u \cdot \partial u / \partial w$

Forward Pass

First, in a forward pass, compute values of all nodes given an input
(The values of each node will be needed during backprop)




Where values computed in the forward pass may be needed

Tricks of the Trade


Regularization

L2 regularization: Introduce a loss term that penalizes the squared magnitude of all parameters. That is, for every weight w in the network, add the term λw^2 to the objective.

Dropout: During training, keep a neuron active with a (keep) probability of p or set it to 0 otherwise.




(a) Standard Neural Net



(b) After applying dropout.


Regularization

Early stopping: Stop training when performance on a validation set has stopped improving




Double descent

Early stopping may not always be an effective strategy considering the “double descent” phenomenon [1]




Learning Rate Schedule

- Observe training losses to understand the effect of different learning rates
- Helpful to decay the learning rate over time. E.g. step decay, exponential decay, etc.



Learning Rate Schedule

- Observe training losses to understand the effect of different learning rates
- Helpful to decay the learning rate over time. E.g. step decay, exponential decay, etc.
- Adaptive learning rate methods like Adagrad, Adam are popular optimizers.



Animation from: <http://cs231n.github.io/neural-networks-3/>

Good reference for optimizers: <https://ruder.io/optimizing-gradient-descent/>



Convolutional Neural Networks

What is in this image?



Windows, porch, steps, door, etc.

Learning Features



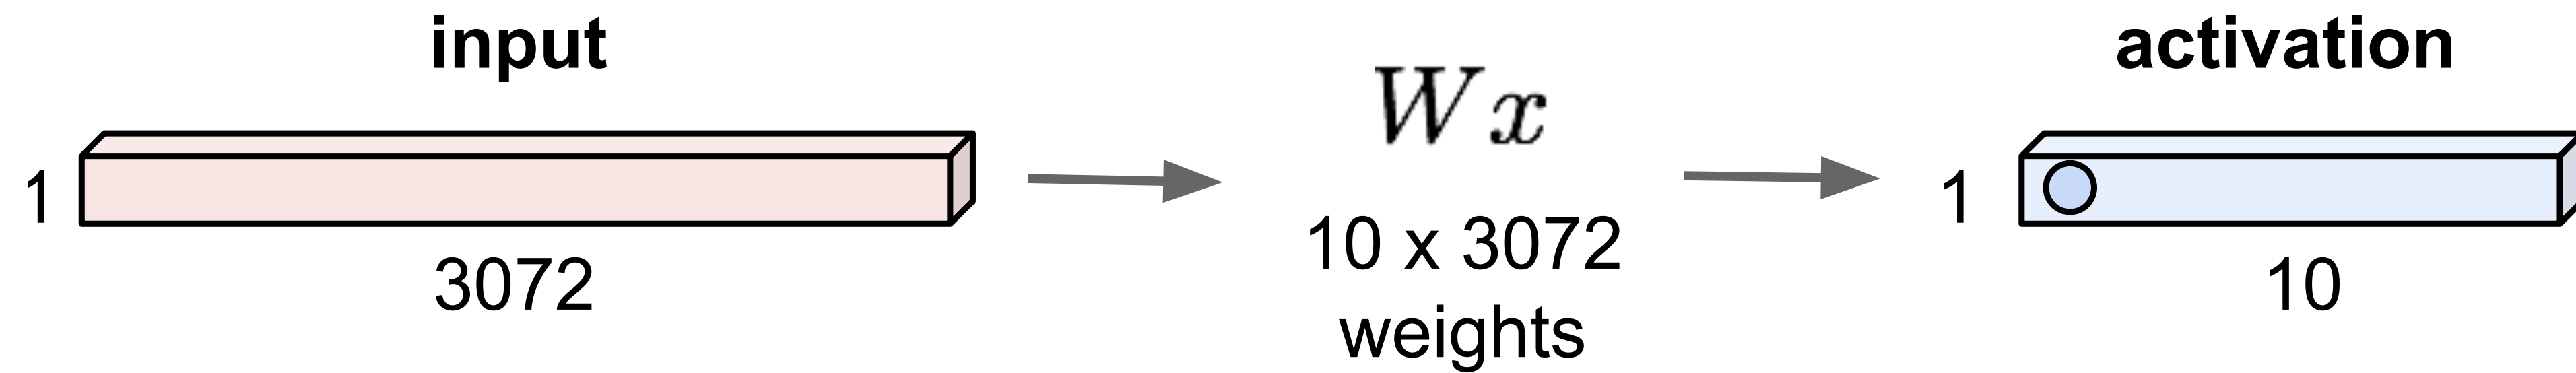
Instead of manually deriving features, can we learn features or representations directly from the data?

Convolutional Neural Networks (CNNs)

- Fully connected (dense) layers have no awareness of spatial information
- Key concept behind convolutional layers is that of ***kernels*** or ***filters***
- Filters slide across an input space to detect spatial patterns (translation invariance) in local regions (locality)


Fully Connected Layers

32x32x3 image -> stretch to 3072 x 1




Convolution Layer


32x32x3 image



5x5x3 filter




Convolution Layer




Convolution Layer



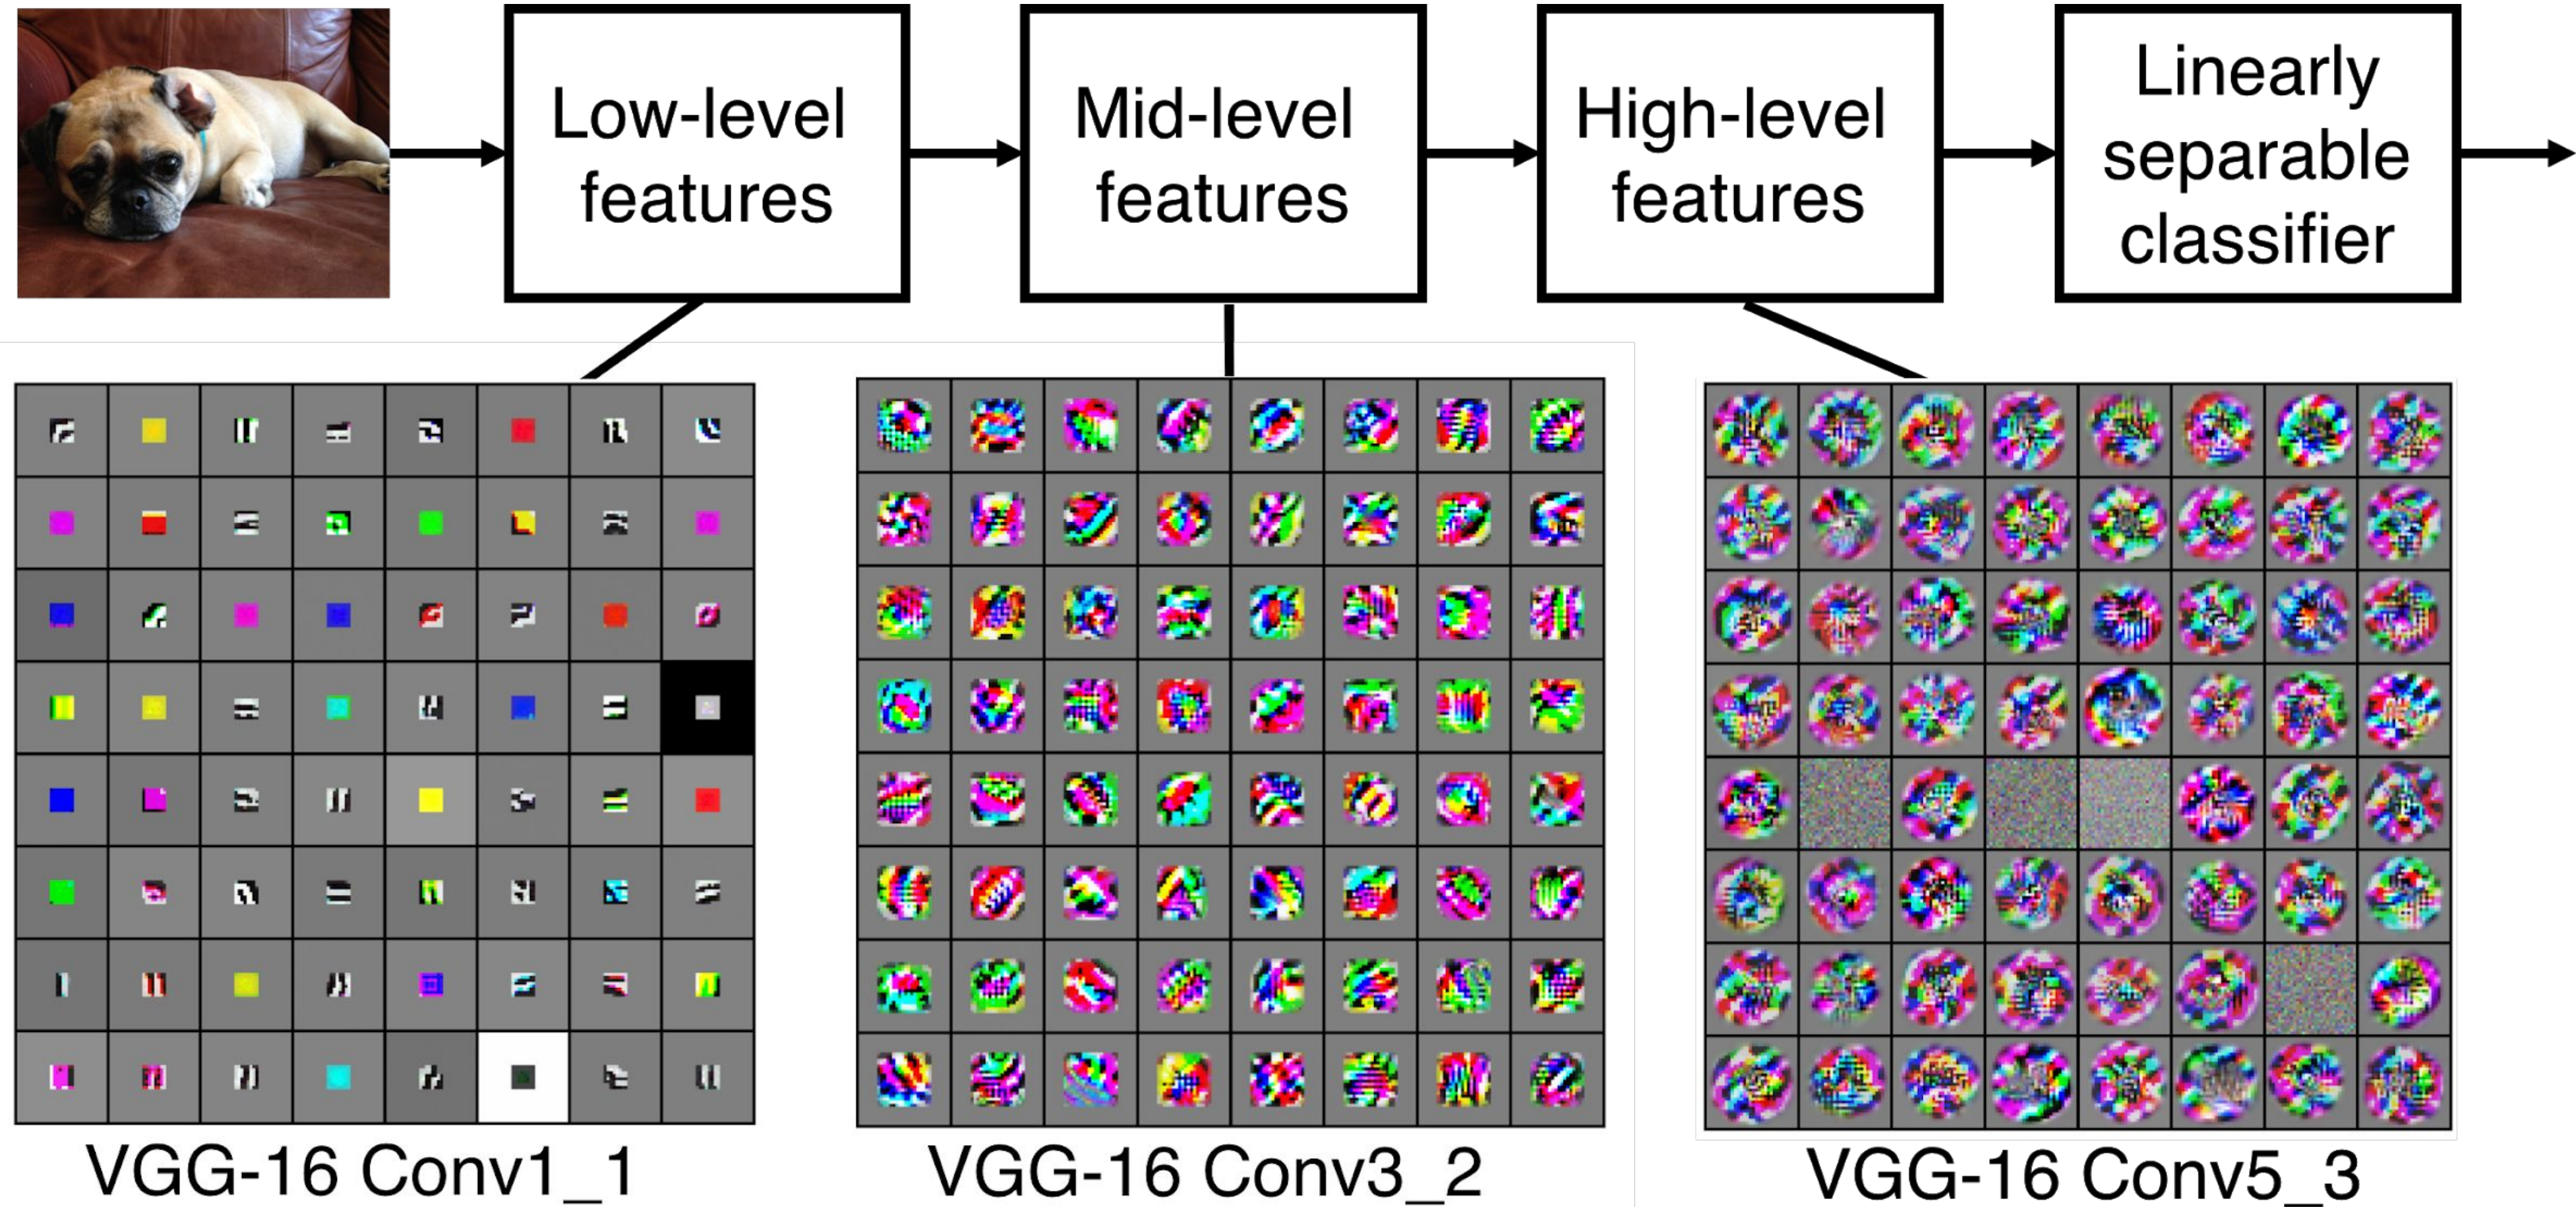
Convolution Layer




Convolutional Neural Network




What do these layers learn?



Convolutional Neural Networks (CNNs)




Stride=1, No padding




Stride=1, Padding, P=1


Convolutional Neural Networks (CNNs)



Stride=1, No padding




Stride=1, Padding, P=1




Stride=1, Padding, P=2


Convolutional Neural Networks (CNNs)




Stride=1, No padding



Stride=1, Padding, P=1




Stride=1, Padding, P=2




Stride=2, No padding


Convolutional Neural Networks (CNNs)




Stride=1, No padding




Stride=1, Padding, P=1



Stride=1, Padding, P=2



Stride=2, No padding




Stride=2, Padding, P=1

Convolution Layers: Summary

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.

Pooling Layer

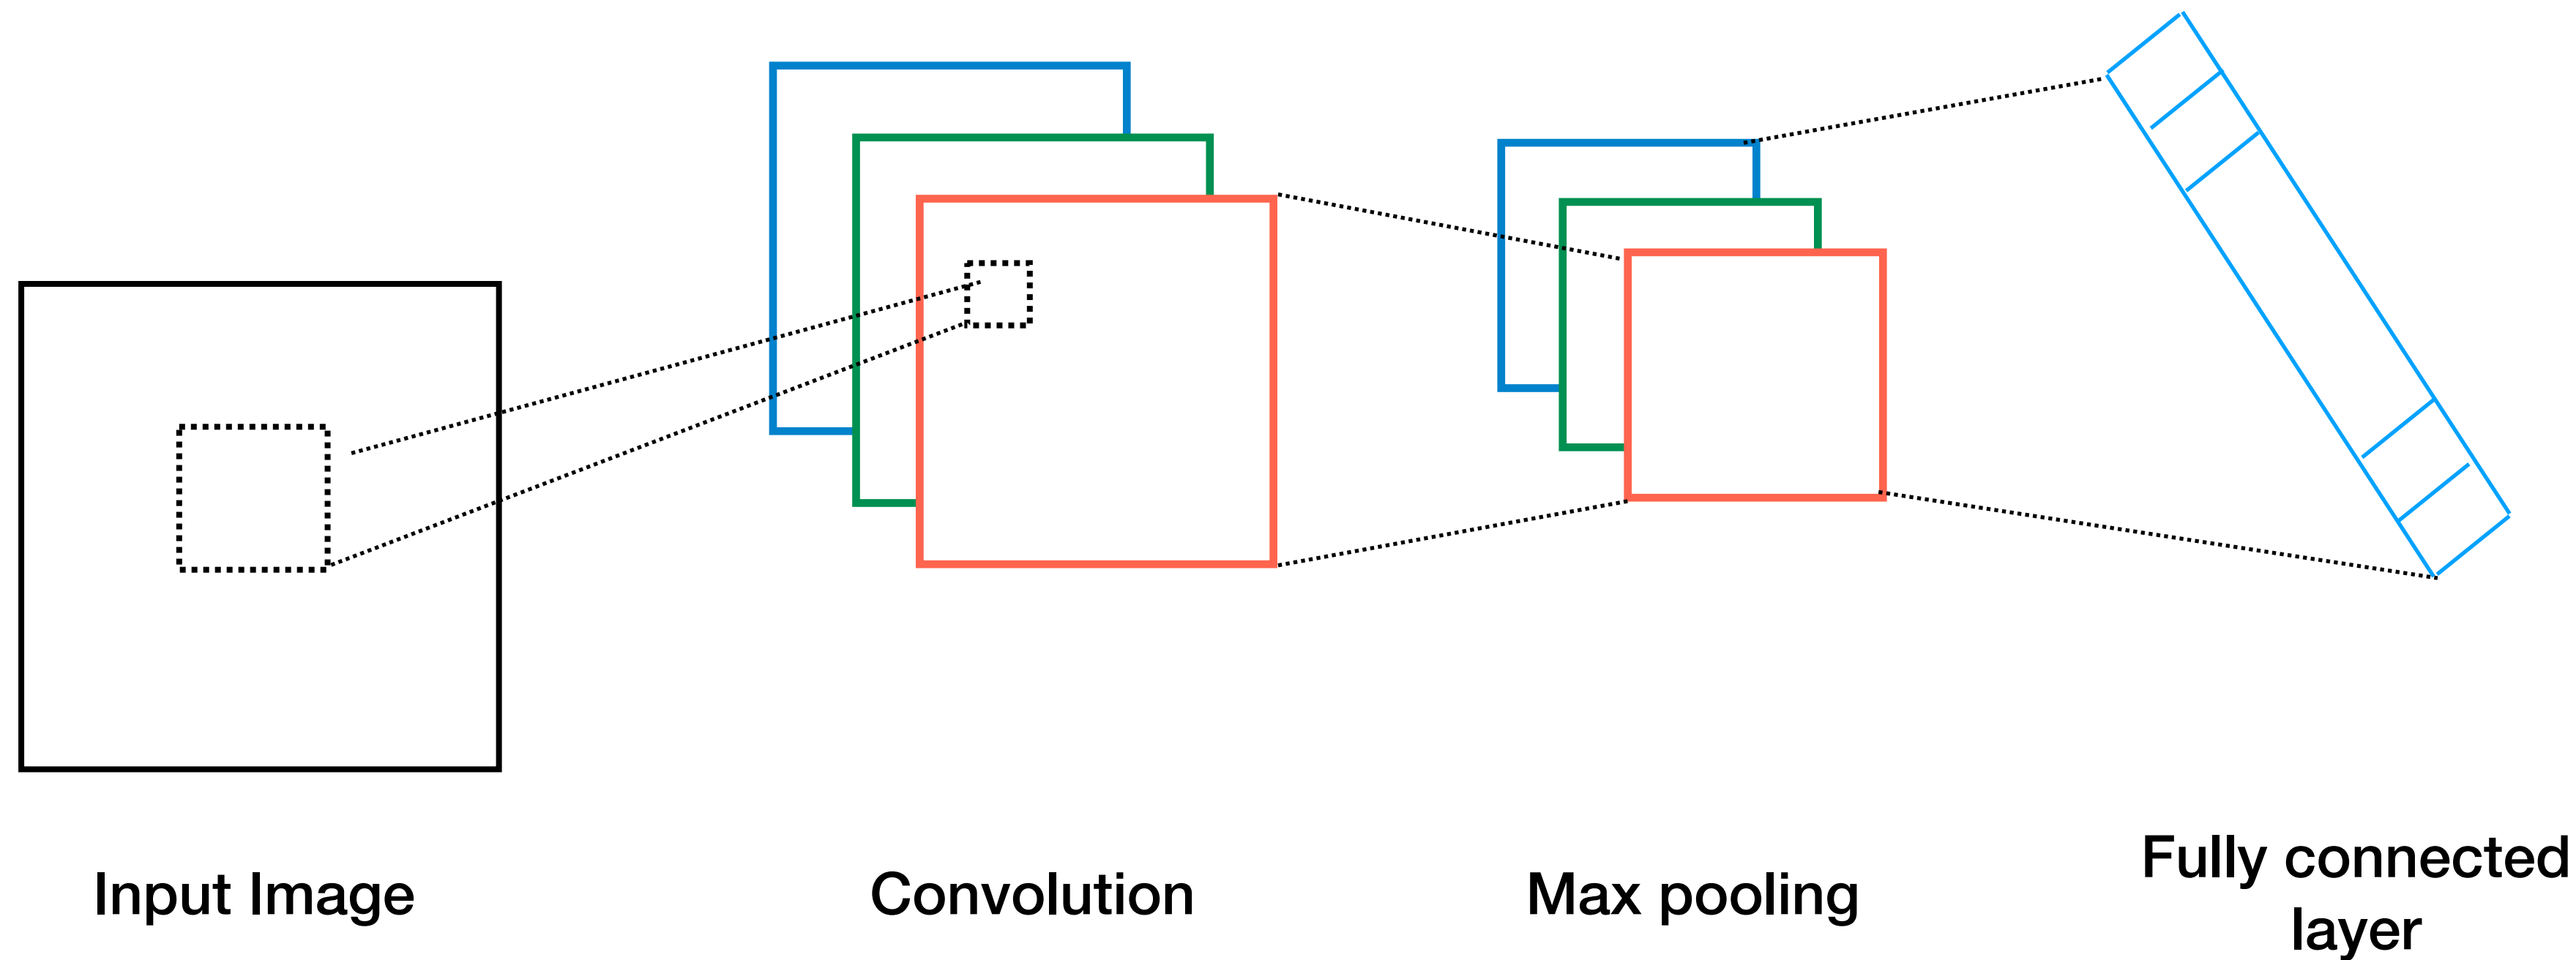


- Why pooling?
Reduce the size of the representation, speed up the computations and make the features a little more robust.
- Max pooling is popularly used in CNNs.

Pooling Layer


- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires two hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$

Convolutional Architectures



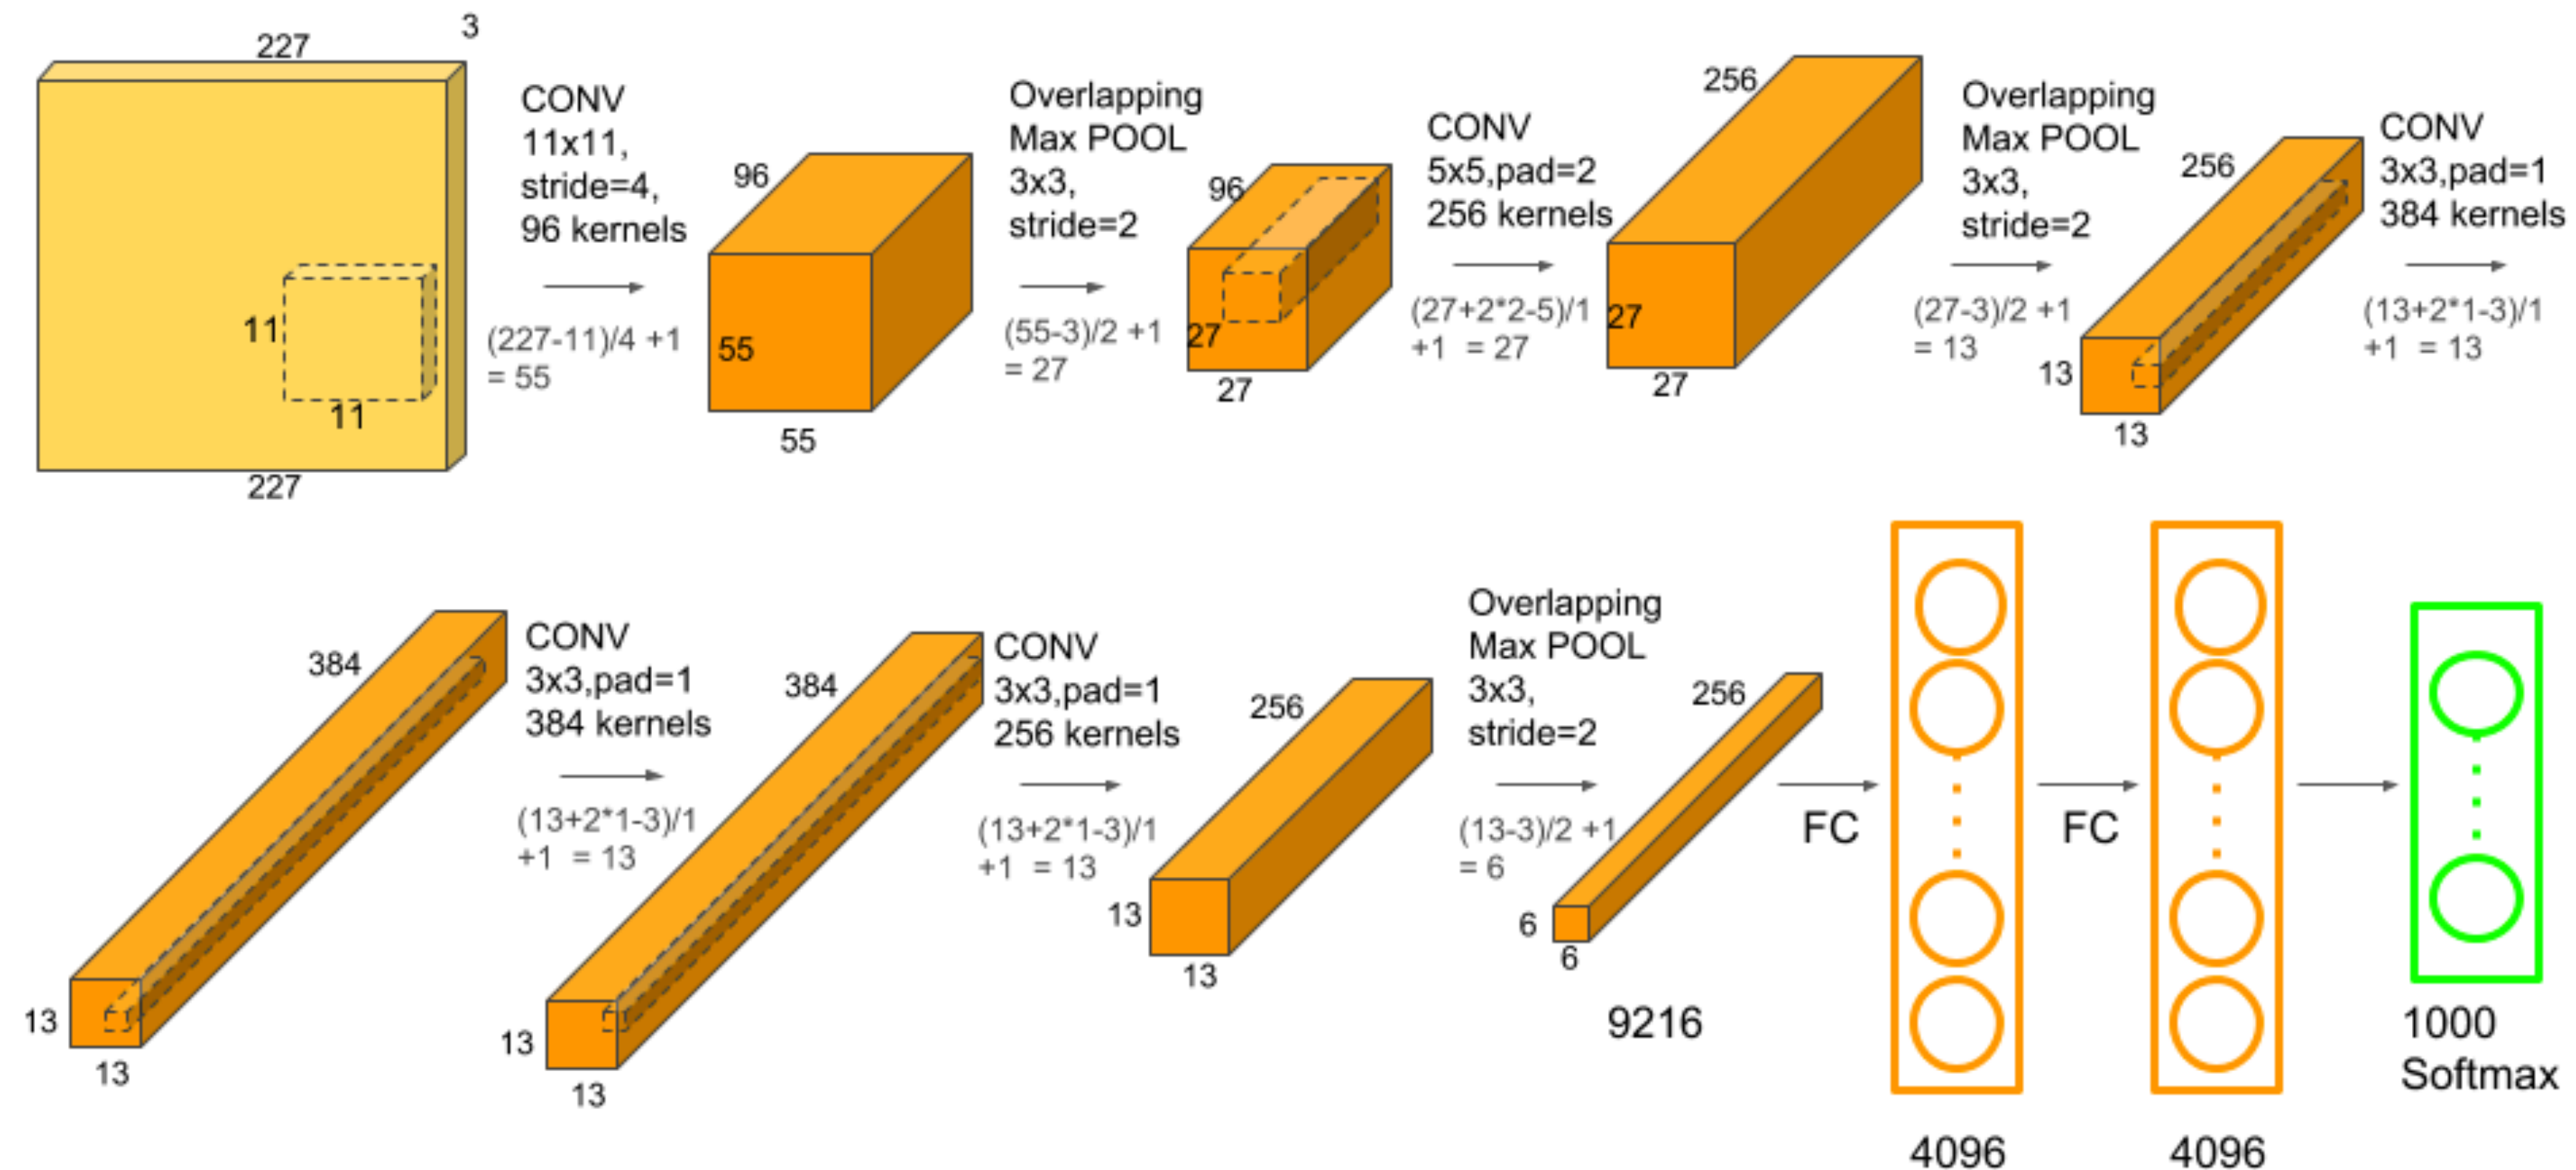
- Block that can be repeated: Convolutional layer, followed by non-linearity (e.g. ReLU) + Max pooling
- Fully connected layers before classification

LeNet-5



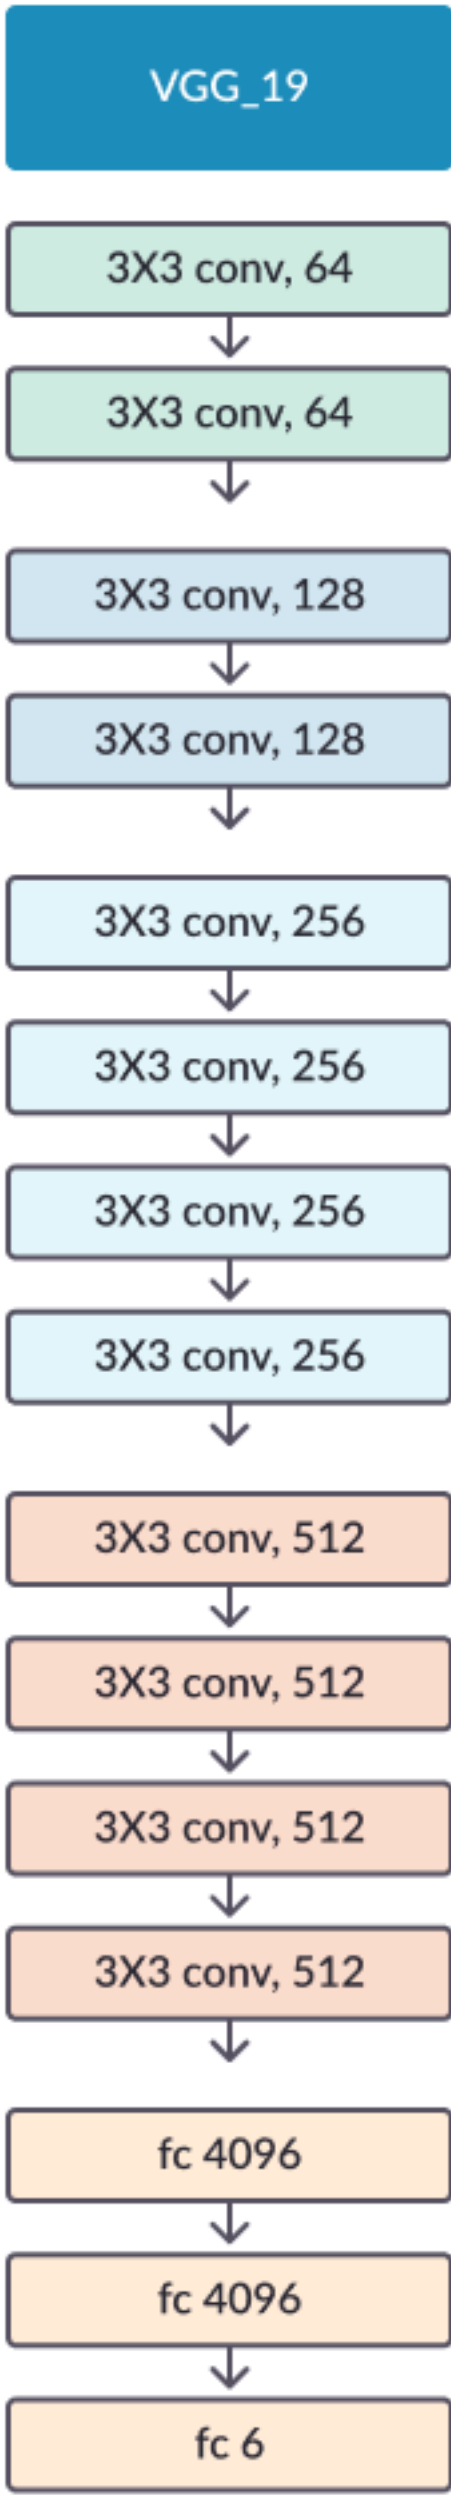
- One of the first successful CNN architectures
- Used to classify images of hand-written digits

AlexNet




- Winner (by a large margin) of the ImageNet challenge in 2012.
- Much larger than previous architectures.


Other Architectures



VGG19



ResNet



VGG: Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." arXiv preprint arXiv:1409.1556 (2014).
ResNet: K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. arXiv preprint arXiv:1512.03385,2015.

Vision and deep neural networks

- Some tasks that have seen a lot of progress:
 - Person/face/object recognition
 - Image segmentation
 - Single person pose estimation
- Some tasks that require a lot more work:
 - Comprehensive scene understanding
 - Reasoning geometrically
 - Generalizing to new domains

NLP/speech and deep neural networks

- Some tasks that have seen a lot of progress:
 - Speech recognition on high-resource languages (e.g. English, Chinese)
 - Low-level NLP tasks like part-of-speech tagging, etc.
- Some tasks that require a lot more work:
 - Natural language understanding
 - NLP/speech technologies for low-resource scenarios
 - Reasoning about large documents